

Prove It First

An AI writing pipeline that isn't allowed to publish a claim it can't back up — and what happened when we actually put that promise to the test.

Someone searches a single word — say, 娛樂城 ("entertainment city," the term for Taiwan's online casino platforms). That's it. No follow-up. But the person who just deposited money and wants to know if it'll come back out, the person who got a referral link from a friend and wants to know if it's real, and the person deciding whether the mobile app is worth their time — they all typed the exact same word, and they all want completely different answers.

THE PROBLEM

One search word hides many real questions

Most content written for a word like this plays it safe and generic — a list of platforms, a bonus percentage, a "top 5" ranking. It answers no one's actual question well, because it's trying to answer everyone's at once.

There's a second problem underneath it. Search engines are increasingly downgrading that kind of generic, safe-sounding content in favor of writing that shows real, first-hand experience — and when the topic touches someone's money, held to an even higher bar of demonstrated trustworthiness. Google's own reviewer guidelines have names for these: **EEAT** (Experience, Expertise, Authoritativeness, Trustworthiness) and **YMYL** ("Your Money or Your Life" — content that can affect what someone does with their money, health, or safety).

Topic 1 is a working answer to both problems at once: figure out the specific hidden question, write like someone who's actually lived it, and — critically — don't let it ship until an independent check confirms it earns the trust it's asking for.

HOW IT WORKS

A three-step relay, with one human checkpoint

01 STEP 1 — AN AI

Figure out what people actually mean

Instead of one generic answer, it produces at least four distinct, specific reasons someone might have searched that word — each with real evidence for why that reason exists and why the usual content ignores it.

02 STEP 2 — A PERSON

Pick which real question to answer

The only non-automated step in the whole process: a human reviews the four-plus candidates and chooses which one this piece of content will actually address.

03 STEP 3 — AN AI

Write like someone who's actually been there

It drafts a first-hand account — specific sensory detail, a genuinely two-sided take, not marketing copy — built around the one question chosen in Step 2.

04 STEP 4 — A SECOND, INDEPENDENT AI

Play the skeptical inspector, then make it prove itself

A separate check reads the draft looking specifically for claims it can't back up. If it finds one, the draft goes back for a rewrite — and gets checked again. This repeats until it passes, or until a cap on attempts is reached.

WHAT ACTUALLY HAPPENED

The inspector catches a real overclaim

This isn't a hypothetical — it's what an actual test run produced. The first draft, describing how smooth the app felt on a few platforms the writer had tried, included one sentence that went further than the evidence supported:

DRAFT — FLAGGED

~~Honestly, these are probably the best connected platforms on the market.~~

REWRITE — PASSED

The steadiest of the few I've personally tried — but I've only tried these few, so I can't say it's the best out there.

The inspector's actual objection: **"best on the market"** is a claim about the entire industry, but the only evidence given was one person's limited experience. It also flagged this as more than a style problem — since it touches how much a reader trusts a platform with their money, an unearned "best" claim is exactly the kind of overreach the EEAT/YMYL bar exists to catch.

THE HONEST RESULT

When we ran it for real, it didn't just pass

The example above is a clean success. But we also ran the full pipeline against a genuinely independent AI model, with nothing staged in advance — and that run is the more informative one, precisely because it didn't go smoothly.

3 / 3

fact-check rounds used, all failed

1

real system bug this run exposed,
since fixed

0

unsupported claims allowed
through anyway

Result: held back, not published

Across all three rounds, the draft got better at admitting its own limits — "I have no professional qualifications in this," "I only tested this once" — but the inspector kept asking for something an honest personal account genuinely can't supply on its own: a named regulator, an actual license number, a source someone could go check. No amount of careful wording substitutes for a real citation. Rather than let softer language slip through as a pass, the system held the line for all three attempts.

That's the system working as intended, not failing. A tool that always finds a way to say "pass" isn't actually checking anything.

What this is worth to a business

- ✓ **A repeatable way to produce content that has to earn trust before it ships** — with a visible, inspectable record of every draft, every objection, and every rewrite, not just a final result to take on faith.
- ✗ **It does not guarantee a pass.** Sometimes the honest answer is "this claim needs evidence we don't actually have" — and the system is built to say that out loud instead of quietly lowering the bar.
- ✓ **That's closer to how a careful editorial or compliance desk works** than to a content mill — it's designed to fail safely, which matters most exactly where the stakes (and the brand risk) are highest.

Derived from Topic 1 of the SEOFUN AI/SEO project. The redline example and the three-round test result above are real system output — not illustrations.