

The In-House Fact-Checker

A private AI, taught with 12 real hand-graded examples, tested honestly against 8 more it never saw — to catch bad SEO advice before it reaches print.

WHAT IT IS

Some published SEO content covers casino sign-up bonuses and second mortgages — “YMYL” topics (Your Money or Your Life) where a wrong claim can cost a reader real money. Sending every article to an outside AI service to check it means the content leaves the building first. This system runs entirely on our own hardware instead, and was taught what a rigorous reviewer looks for using a small, real, hand-graded example set.

HOW IT WAS TESTED, HONESTLY

- **12 real examples** taught it, each scored by hand with the specific reasons why.
- **8 more real examples were held back** and never shown to it while learning — the fair test.
- The **untrained** and **newly-taught** versions were run on the same 8 unseen examples, side by side.

A REAL ONE IT CAUGHT

BEFORE TRAINING

On a false claim it had never seen (“repeat your keyword 20× in the H1 tag to rank higher”): “The text asserts an SEO tactic without providing evidence or credible sources.”

AFTER TRAINING

Same never-seen article: “Google has repeatedly and publicly stated that keyword repetition in H1 tags is a ranking penalty, not a ranking factor.”

THE HONEST NUMBERS (0-100 scale)

Genuinely good content	64 → 88	scored higher
Overpromising / risky claims	20 → 5	scored lower
Fabricated SEO “theories”	13 → 2	scored lower, more consistently
Obvious keyword-stuffing	20 → 10	scored lower

One honest catch, found the same careful way: the **untrained** version scored the identical, unchanged article **60/100 and 85/100** on separate tries — a real self-contradiction, not sampling noise. Strict JSON formatting was followed **43 of 48 times (89.6%)**, not every time. Tested on 8 held-out examples — directionally consistent, not a large-sample result. Both facts are reported in full, not smoothed over.