

The Negotiating AI

Two AIs with opposite goals, forced into a structured back-and-forth until they reach an honest compromise — with the full argument kept on record.

WHAT IT IS

Content teams live with a real tension: the person chasing clicks and the person keeping things legal and honest rarely want the same headline. Usually one side just wins, quietly, with no record of why. This system simulates the negotiation instead — a Planner AI built to maximize clicks, an Auditor AI built to enforce the real rules, forced to work it out through structured rounds rather than one side overruling the other.

HOW IT WORKS

Start -> Planner drafts -> Auditor checks -> Decision

Decision -> passed? -> done

Decision -> rejected? -> back to Planner (with the specific reason why)

A REAL ONE IT CAUGHT

ROUND 0 — REJECTED

Planner's headline promised “guaranteed profit” and “guaranteed payout.” Auditor named both banned phrases specifically and explained why each is a problem.

ROUND 1 — PASSED

Planner reframed around verifiable licensing, identity checks, and real withdrawal processing times — still built to be clicked, just no longer an unconditional promise.

THE PART WE DIDN'T EXPECT

The company's compliance manual was written for mortgage content. Asked to review a casino ad instead, the Auditor correctly said those specific rules didn't apply here — and reasoned from the actual banned phrases instead of forcing an irrelevant citation. That's real judgment, not a rigid checklist.

Being direct about what this run does and doesn't prove: this is **one real test scenario**, not a large-scale trial. The negotiation converged in 2 rounds with no formatting errors — a clean result — but the “walk away without agreeing” safety valve, and how the Auditor behaves when the compliance manual's rules genuinely do apply, are both still untested against a live run.